

# Assignment 1

---

## Exercise 1 - Reflections on GPU-accelerated Computing

1

- CPUs are faster and better at handling complex instructions
- GPUs have a lot more cores which means they can work on a lot of tasks in parallel
- CPUs usually have higher clock speeds and therefore can execute a single instruction faster than a GPU

2

Rank	Name	GPU model	Rpeak / Power (TFlops/kW)
1	Frontier	AMD Instinct MI250X	73.95
2	Fugaku	Fujitsu A64FX	14.78
3	LUMI	AMD Instinct MI250X	51.36
4	Leonardo	Nvidia Ampere A100	32.14
5	Summit	Nvidia Tesla V100 / Volta GV100	14.66
6	Sierra	Nvidia Tesla V100 / Volta GV100	12.64
7	Sunway TaihuLight	-	6.05
8	Perlmutter	Nvidia Ampere A100	27.04
9	Selene	Nvidia Ampere A100	23.81
10	Tianhe-2A	Matrix-2000	3.30

- 9 out of 10 have a GPU
- Out of those 5 are by Nvidia 2 by AMD and one each from Fujitsu and Matrix
- source: <https://www.top500.org/lists/top500/2023/06/>

## Exercise 2 - Query Nvidia GPU Compute Capability

```

✓ 1s !./deviceQuery
./deviceQuery Starting...

  CUDA Device Query (Runtime API) version (CUDA static linking)

Detected 1 CUDA Capable device(s)

Device 0: "Tesla T4"
  CUDA Driver Version / Runtime Version      12.0 / 11.8
  CUDA Capability Major/Minor version number: 7.5
  Total amount of global memory:             15102 MBytes (15835398144 bytes)
  (040) Multiprocessors, (064) CUDA Cores/MP: 2560 CUDA Cores
  GPU Max Clock rate:                       1590 MHz (1.59 GHz)
  Memory Clock rate:                        5001 Mhz
  Memory Bus Width:                         256-bit
  L2 Cache Size:                           4194304 bytes
  Maximum Texture Dimension Size (x,y,z)     1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
  Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
  Total amount of constant memory:           65536 bytes
  Total amount of shared memory per block:    49152 bytes
  Total shared memory per multiprocessor:     65536 bytes
  Total number of registers available per block: 65536
  Warp size:                                32
  Maximum number of threads per multiprocessor: 1024
  Maximum number of threads per block:        1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size (x,y,z):  (2147483647, 65535, 65535)
  Maximum memory pitch:                      2147483647 bytes
  Texture alignment:                         512 bytes
  Concurrent copy and kernel execution:       Yes with 3 copy engine(s)
  Run time limit on kernels:                  No
  Integrated GPU sharing Host Memory:         No
  Support host page-locked memory mapping:    Yes
  Alignment requirement for Surfaces:         Yes
  Device has ECC support:                     Enabled
  Device supports Unified Addressing (UVA):   Yes
  Device supports Managed Memory:             Yes
  Device supports Compute Preemption:         Yes
  Supports Cooperative Kernel Launch:         Yes
  Supports MultiDevice Co-op Kernel Launch:   Yes
  Device PCI Domain ID / Bus ID / location ID: 0 / 0 / 4
  Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 12.0, CUDA Runtime Version = 11.8, NumDevs = 1
Result = PASS

```

The compute capability is 7.5.

- Memory Clock rate: 5001 Mhz
- Bus width: 256-bit
- DDR: 2
- Memory bandwidth:  $5001 * 256 / 8 / 1024 * 2 = 312 \text{ GB/s}$

! ./bandwidthTest

[CUDA Bandwidth Test] - Starting...  
Running on...

Device 0: Tesla T4  
Quick Mode

Host to Device Bandwidth, 1 Device(s)

PINNED Memory Transfers

Transfer Size (Bytes)	Bandwidth(GB/s)
32000000	10.1

Device to Host Bandwidth, 1 Device(s)

PINNED Memory Transfers

Transfer Size (Bytes)	Bandwidth(GB/s)
32000000	11.2

Device to Device Bandwidth, 1 Device(s)

PINNED Memory Transfers

Transfer Size (Bytes)	Bandwidth(GB/s)
32000000	239.4

Result = PASS

NOTE: The CUDA Samples are not meant for performance measurements. Results may vary when GPU Boost is enabled.

The bandwidth from the test is lower than the value calculated.

## Exercise 3 - Rodinia CUDA benchmarks and Comparison with CPU

I've carried out experiments the heartwall and k-means tests. Changing the Makefile was not necessary however I removed the unnecessary file compilations to simplify the process and shorten the compilation times. The two selected tasks are imaging and data mining tasks which can be done in parallel. Looking at the results in the below figures it can be seen that for both tests the execution time is more than twice as fast on a GPU compared to a CPU. While both executions can be parallelised the OMP code has been run using a single thread while the CUDA code uses 256 threads per block for the imaging task. This means that it can do many more calculations at the same time.

```
✓ [109] ! time ./heartwall ../../../../data/heartwall/test.avi 10
0s

WG size of kernel = 256
frame progress: 0 1 2 3 4 5 6 7 8 9

real    0m0.435s
user    0m0.136s
sys     0m0.249s

✓ [110] ! nvprof ./heartwall ../../../../data/heartwall/test.avi 10
0s

WG size of kernel = 256
==45201== NVPROF is profiling process 45201, command: ./heartwall ../../../../data/heartwall/test.avi 10
frame progress: 0 1 2 3 4 5 6 7 8 9
==45201== Profiling application: ./heartwall ../../../../data/heartwall/test.avi 10
==45201== Profiling result:
   Type  Time(%)   Time    Calls   Avg      Min      Max  Name
GPU activities:  97.96% 122.25ms    10 12.225ms 13.024us 13.590ms kernel(void)
                2.03%  2.5394ms    26  97.669us 1.3120us 273.37us [CUDA memcpy HtoD]
                0.01%  10.784us     4  2.6960us 2.5920us  2.8160us [CUDA memcpy DtoH]
API calls:      70.43% 231.58ms   623 371.73us 2.0820us 228.13ms cudaMalloc
                27.93%  91.827ms    18  5.1015ms 4.2070us 13.590ms cudaMemcpy
                1.30%  4.2592ms   623  6.8360us 2.6220us 148.27us cudaFree
                0.24%  792.45us    12 66.037us 10.028us 81.935us cudaMemcpyToSymbol
                0.05%  177.59us    10 17.759us 13.920us 34.008us cudaLaunchKernel
                0.04%  121.13us   101  1.1990us   133ns 49.039us cuDeviceGetAttribute
                0.01%  24.702us     1 24.702us 24.702us 24.702us cuDeviceGetName
                0.00%  6.4520us     1  6.4520us 6.4520us 6.4520us cuDeviceGetPCIBusId
                0.00%  2.6080us     2  1.3040us   165ns 2.4430us cuDeviceGet
                0.00%  1.9790us     3    659ns   224ns 1.5300us cuDeviceGetCount
                0.00%    495ns     1    495ns   495ns   495ns cuModuleGetLoadingMode
                0.00%    471ns     1    471ns   471ns   471ns cuDeviceTotalMem
                0.00%    232ns     1    232ns   232ns   232ns cuDeviceGetUuid

✓ [111] %cd /content/drive/MyDrive/DD2360/rodinia_3.1/bin/linux/omp
0s

/content/drive/MyDrive/DD2360/rodinia_3.1/bin/linux/omp

✓ [64] ! chmod +x heartwall
0s

✓ [112] ! time ./heartwall ../../../../data/heartwall/test.avi 10 1
11s

num of threads: 1
frame progress: 0 1 2 3 4 5 6 7 8 9

real    0m11.396s
user    0m11.193s
sys     0m0.051s
```

```
✓ [102] ! time ./kmeans -i /content/drive/MyDrive/DD2360/rod/rodinia_3.1/data/kmeans/819200.txt
```

```
I/O completed
```

```
Number of objects: 819200
Number of features: 34
Iterated 2 times
Number of Iteration: 1
```

```
real    0m2.108s
user    0m1.648s
sys     0m0.407s
```

```
✓ [107] ! nvprof ./kmeans -i /content/drive/MyDrive/DD2360/rod/rodinia_3.1/data/kmeans/819200.txt
```

```
==44770== NvPROF is profiling process 44770, command: ./kmeans -i /content/drive/MyDrive/DD2360/rod/rodinia_3.1/data/kmeans/819200.txt
```

```
I/O completed
```

```
Number of objects: 819200
Number of features: 34
Iterated 2 times
Number of Iteration: 1
```

```
==44770== Profiling application: ./kmeans -i /content/drive/MyDrive/DD2360/rod/rodinia_3.1/data/kmeans/819200.txt
==44770== Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	69.40%	25.318ms	5	5.0637ms	1.3750us	24.247ms	[CUDA memcpy HtoD]
	15.42%	5.6269ms	1	5.6269ms	5.6269ms	5.6269ms	invert_mapping(float*, float*, int, int)
	12.40%	4.5230ms	2	2.2615ms	2.2563ms	2.2668ms	kmeansPoint(float*, int, int, int, int*, float*, float*, int*)
	2.77%	1.0115ms	2	505.74us	408.25us	603.23us	[CUDA memcpy DtoH]
API calls:	84.06%	212.60ms	4	53.149ms	72.386us	212.22ms	cudaMalloc
	12.23%	30.929ms	7	4.4185ms	71.935us	24.445ms	cudaMemcpy
	1.79%	4.5301ms	2	2.2651ms	2.2598ms	2.2703ms	cudaThreadSynchronize
	1.36%	3.4268ms	4	856.69us	231.07us	1.1591ms	cudaFree
	0.45%	1.1321ms	1	1.1321ms	1.1321ms	1.1321ms	cuDeviceGetPCIBusId
	0.05%	116.79us	101	1.1560us	149ns	48.870us	cuDeviceGetAttribute
	0.04%	93.515us	3	31.171us	28.102us	34.190us	cudaLaunchKernel
	0.01%	30.113us	6	5.0180us	1.3350us	17.137us	cudaBindTexture
	0.01%	25.208us	1	25.208us	25.208us	25.208us	cuDeviceGetName
	0.00%	6.9870us	3	2.3290us	2.0910us	2.6780us	cudaSetDevice
	0.00%	3.1930us	2	1.5960us	1.4670us	1.7260us	cudaMemcpyToSymbol
	0.00%	1.4550us	3	485ns	195ns	1.0340us	cuDeviceGetCount
	0.00%	1.2140us	6	202ns	117ns	378ns	cudaCreateChannelDesc
	0.00%	937ns	2	468ns	170ns	767ns	cuDeviceGet
	0.00%	604ns	1	604ns	604ns	604ns	cuModuleGetLoadingMode
	0.00%	377ns	1	377ns	377ns	377ns	cuDeviceTotalMem
	0.00%	247ns	1	247ns	247ns	247ns	cuDeviceGetUuid

```
✓ [103] %cd /content/drive/MyDrive/DD2360/rod/rodinia_3.1/bin/linux/omp
```

```
/content/drive/MyDrive/DD2360/rod/rodinia_3.1/bin/linux/omp
```

```
✓ [89] !chmod +x kmeans
```

```
[104] ! time ./kmeans -i /content/drive/MyDrive/DD2360/rod/rodinia_3.1/data/kmeans/819200.txt
```

```
I/O completed
```

```
num of threads = 1
number of Clusters 5
number of Attributes 34
```

```
Time for process: 3.201367
```

```
real    0m4.910s
user    0m4.661s
sys     0m0.183s
```

## Exercise 4 - Run a HelloWorld on AMD GPU

To launch the code on the AMD GPUs in Dardel. First, one needs to get a GPU allocation. Then the executable created can be launched on the GPU using the **srun** command and specifying which node the program should run on with the **-n** flag.

```
[aorucu@uan01:~/Private/hw1> srun -n 1 ./HelloWorld
```

```
System minor 0
```

```
System major 9
```

```
agent prop name
```

```
input string:
```

```
GdkknVnqkc
```

```
output string:
```

```
HelloWorld
```

```
Passed!
```

```
[aorucu@uan01:~/Private/hw1> ]
```