

# Assignment 1

---

## Exercise 1 - Reflections on GPU-accelerated Computing

1

- CPUs are faster and better at handling complex instructions
- GPUs have a lot more cores which means they can work on a lot of tasks in parallel
- 

2

Rank	Name	GPU model	Rpeak / Power (TFlops/kW)
1	Frontier	AMD Instinct MI250X	73.95
2	Fugaku	Fujitsu A64FX	14.78
3	LUMI	AMD Instinct MI250X	51.36
4	Leonardo	Nvidia Ampere A100	32.14
5	Summit	Nvidia Tesla V100 / Volta GV100	14.66
6	Sierra	Nvidia Tesla V100 / Volta GV100	12.64
7	Sunway TaihuLight	-	6.05
8	Perlmutter	Nvidia Ampere A100	27.04
9	Selene	Nvidia Ampere A100	23.81
10	Tianhe-2A	Matrix-2000	3.30

- 9 out of 10 have a GPU
- Out of those 5 are by Nvidia 2 by AMD and one each from Fujitsu and Matrix
- source: <https://www.top500.org/lists/top500/2023/06/>

## Exercise 2 - Query Nvidia GPU Compute Capability

```

✓ 1s !./deviceQuery
./deviceQuery Starting...

    CUDA Device Query (Runtime API) version (CUDA static linking)

Detected 1 CUDA Capable device(s)

Device 0: "Tesla T4"
  CUDA Driver Version / Runtime Version      12.0 / 11.8
  CUDA Capability Major/Minor version number: 7.5
  Total amount of global memory:             15102 MBytes (15835398144 bytes)
  (040) Multiprocessors, (064) CUDA Cores/MP: 2560 CUDA Cores
  GPU Max Clock rate:                       1590 MHz (1.59 GHz)
  Memory Clock rate:                        5001 Mhz
  Memory Bus Width:                         256-bit
  L2 Cache Size:                           4194304 bytes
  Maximum Texture Dimension Size (x,y,z)     1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
  Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
  Total amount of constant memory:            65536 bytes
  Total amount of shared memory per block:    49152 bytes
  Total shared memory per multiprocessor:     65536 bytes
  Total number of registers available per block: 65536
  Warp size:                                32
  Maximum number of threads per multiprocessor: 1024
  Maximum number of threads per block:        1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size (x,y,z):  (2147483647, 65535, 65535)
  Maximum memory pitch:                      2147483647 bytes
  Texture alignment:                         512 bytes
  Concurrent copy and kernel execution:       Yes with 3 copy engine(s)
  Run time limit on kernels:                  No
  Integrated GPU sharing Host Memory:         No
  Support host page-locked memory mapping:    Yes
  Alignment requirement for Surfaces:         Yes
  Device has ECC support:                     Enabled
  Device supports Unified Addressing (UVA):    Yes
  Device supports Managed Memory:             Yes
  Device supports Compute Preemption:         Yes
  Supports Cooperative Kernel Launch:         Yes
  Supports MultiDevice Co-op Kernel Launch:   Yes
  Device PCI Domain ID / Bus ID / location ID: 0 / 0 / 4
  Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 12.0, CUDA Runtime Version = 11.8, NumDevs = 1
Result = PASS

```

The compute capability is 7.5.

- Memory Clock rate: 5001 Mhz
- Bus width: 256-bit
- DDR: 2
- Memory bandwidth:  $5001 * 256 / 8 / 1024 * 2 = 312 \text{ GB/s}$

! ./bandwidthTest

[CUDA Bandwidth Test] - Starting...  
Running on...

Device 0: Tesla T4  
Quick Mode

Host to Device Bandwidth, 1 Device(s)

PINNED Memory Transfers

Transfer Size (Bytes)	Bandwidth(GB/s)
32000000	10.1

Device to Host Bandwidth, 1 Device(s)

PINNED Memory Transfers

Transfer Size (Bytes)	Bandwidth(GB/s)
32000000	11.2

Device to Device Bandwidth, 1 Device(s)

PINNED Memory Transfers

Transfer Size (Bytes)	Bandwidth(GB/s)
32000000	239.4

Result = PASS

NOTE: The CUDA Samples are not meant for performance measurements. Results may vary when GPU Boost is enabled.

The bandwidth from the test is lower than the value calculated.

## Exercise 3 - Rodinia CUDA benchmarks and Comparison with CPU

In both images provided the results on the top are on a GPU and on the bottom on one thread using OpenMP. The programs executed on CUDA is faster than both of the programs - heartwall and k-means. These tasks are imaging and data mining tasks which can be done in parallel. For this reason GPU performs it much faster.

```
✓ [109] ! time ./heartwall ../../../../data/heartwall/test.avi 10
0s

WG size of kernel = 256
frame progress: 0 1 2 3 4 5 6 7 8 9

real    0m0.435s
user    0m0.136s
sys     0m0.249s

✓ [110] ! nvprof ./heartwall ../../../../data/heartwall/test.avi 10
0s

WG size of kernel = 256
==45201== NVPROF is profiling process 45201, command: ./heartwall ../../../../data/heartwall/test.avi 10
frame progress: 0 1 2 3 4 5 6 7 8 9
==45201== Profiling application: ./heartwall ../../../../data/heartwall/test.avi 10
==45201== Profiling result:
   Type  Time(%)   Time    Calls   Avg      Min      Max  Name
GPU activities:  97.96% 122.25ms    10 12.225ms 13.024us 13.590ms kernel(void)
                2.03%  2.5394ms    26  97.669us 1.3120us 273.37us [CUDA memcpy HtoD]
                0.01%  10.784us     4  2.6960us 2.5920us  2.8160us [CUDA memcpy DtoH]
API calls:      70.43% 231.58ms   623 371.73us 2.0820us 228.13ms cudaMalloc
                27.93%  91.827ms    18  5.1015ms 4.2070us 13.590ms cudaMemcpy
                1.30%  4.2592ms   623  6.8360us 2.6220us 148.27us cudaFree
                0.24%  792.45us    12 66.037us 10.028us  81.935us cudaMemcpyToSymbol
                0.05%  177.59us    10 17.759us 13.920us  34.008us cudaLaunchKernel
                0.04%  121.13us   101  1.1990us   133ns  49.039us cuDeviceGetAttribute
                0.01%  24.702us     1 24.702us 24.702us 24.702us cuDeviceGetName
                0.00%  6.4520us     1  6.4520us  6.4520us  6.4520us cuDeviceGetPCIBusId
                0.00%  2.6080us     2  1.3040us   165ns  2.4430us cuDeviceGet
                0.00%  1.9790us     3    659ns   224ns  1.5300us cuDeviceGetCount
                0.00%    495ns     1    495ns   495ns    495ns cuModuleGetLoadingMode
                0.00%    471ns     1    471ns   471ns    471ns cuDeviceTotalMem
                0.00%    232ns     1    232ns   232ns    232ns cuDeviceGetUuid

✓ [111] %cd /content/drive/MyDrive/DD2360/rodinia_3.1/bin/linux/omp
0s

/content/drive/MyDrive/DD2360/rodinia_3.1/bin/linux/omp

✓ [64] ! chmod +x heartwall
0s

✓ [112] ! time ./heartwall ../../../../data/heartwall/test.avi 10 1
11s

num of threads: 1
frame progress: 0 1 2 3 4 5 6 7 8 9

real    0m11.396s
user    0m11.193s
sys     0m0.051s
```

```
✓ [102] ! time ./kmeans -i /content/drive/MyDrive/DD2360/rod/rodinia_3.1/data/kmeans/819200.txt
```

```
I/O completed
```

```
Number of objects: 819200
Number of features: 34
Iterated 2 times
Number of Iteration: 1
```

```
real    0m2.108s
user    0m1.648s
sys     0m0.407s
```

```
✓ [107] ! nvprof ./kmeans -i /content/drive/MyDrive/DD2360/rod/rodinia_3.1/data/kmeans/819200.txt
```

```
==44770== NvPROF is profiling process 44770, command: ./kmeans -i /content/drive/MyDrive/DD2360/rod/rodinia_3.1/data/kmeans/819200.txt
```

```
I/O completed
```

```
Number of objects: 819200
Number of features: 34
Iterated 2 times
Number of Iteration: 1
```

```
==44770== Profiling application: ./kmeans -i /content/drive/MyDrive/DD2360/rod/rodinia_3.1/data/kmeans/819200.txt
==44770== Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	69.40%	25.318ms	5	5.0637ms	1.3750us	24.247ms	[CUDA memcpy HtoD]
	15.42%	5.6269ms	1	5.6269ms	5.6269ms	5.6269ms	invert_mapping(float*, float*, int, int)
	12.40%	4.5230ms	2	2.2615ms	2.2563ms	2.2668ms	kmeansPoint(float*, int, int, int, int*, float*, float*, int*)
	2.77%	1.0115ms	2	505.74us	408.25us	603.23us	[CUDA memcpy DtoH]
API calls:	84.06%	212.60ms	4	53.149ms	72.386us	212.22ms	cudaMalloc
	12.23%	30.929ms	7	4.4185ms	71.935us	24.445ms	cudaMemcpy
	1.79%	4.5301ms	2	2.2651ms	2.2598ms	2.2703ms	cudaThreadSynchronize
	1.36%	3.4268ms	4	856.69us	231.07us	1.1591ms	cudaFree
	0.45%	1.1321ms	1	1.1321ms	1.1321ms	1.1321ms	cuDeviceGetPCIBusId
	0.05%	116.79us	101	1.1560us	149ns	48.870us	cuDeviceGetAttribute
	0.04%	93.515us	3	31.171us	28.102us	34.190us	cudaLaunchKernel
	0.01%	30.113us	6	5.0180us	1.3350us	17.137us	cudaBindTexture
	0.01%	25.208us	1	25.208us	25.208us	25.208us	cuDeviceGetName
	0.00%	6.9870us	3	2.3290us	2.0910us	2.6780us	cudaSetDevice
	0.00%	3.1930us	2	1.5960us	1.4670us	1.7260us	cudaMemcpyToSymbol
	0.00%	1.4550us	3	485ns	195ns	1.0340us	cuDeviceGetCount
	0.00%	1.2140us	6	202ns	117ns	378ns	cudaCreateChannelDesc
	0.00%	937ns	2	468ns	170ns	767ns	cuDeviceGet
	0.00%	604ns	1	604ns	604ns	604ns	cuModuleGetLoadingMode
	0.00%	377ns	1	377ns	377ns	377ns	cuDeviceTotalMem
	0.00%	247ns	1	247ns	247ns	247ns	cuDeviceGetUuid

```
✓ [103] %cd /content/drive/MyDrive/DD2360/rod/rodinia_3.1/bin/linux/omp
```

```
/content/drive/MyDrive/DD2360/rod/rodinia_3.1/bin/linux/omp
```

```
✓ [89] !chmod +x kmeans
```

```
[104] ! time ./kmeans -i /content/drive/MyDrive/DD2360/rod/rodinia_3.1/data/kmeans/819200.txt
```

```
I/O completed
```

```
num of threads = 1
number of Clusters 5
number of Attributes 34
```

```
Time for process: 3.201367
```

```
real    0m4.910s
user    0m4.661s
sys     0m0.183s
```

## Exercise 4 - Run a HelloWorld on AMD GPU

To launch the code on the AMD GPUs in Dardel. First, one needs to get a GPU allocation. Then the executable created can be launched on the GPU using the **srun** command and specifying which node the program should run on with the **-n** flag.

```
[aorucu@uan01:~/Private/hw1> srun -n 1 ./HelloWorld
```

```
System minor 0
```

```
System major 9
```

```
agent prop name
```

```
input string:
```

```
GdkknVnqkc
```

```
output string:
```

```
HelloWorld
```

```
Passed!
```

```
[aorucu@uan01:~/Private/hw1> ]
```